

January 18, 2020 Notes:

- Guided Tour of Chapter 1
- Before class, read chapter in book corresponding to lecture
- Markov Processes and Markov Reward Processes
 - ↳ This lecture is considered basic
 - ↳ Get comfortable with mathematical foundations and foundational code.
 - ↳ Markov Process is often referred to as a Markov Chain
 - ↳ Process: Time-sequence random outcomes
 - ↳ State: Internal Representation S_t
 - ↳ We are interested in $P[S_{t+1} | S_t, S_{t-1}, \dots, S_0]$
- Consider random walks of stock prices $X_t = S_t$
$$P[X_{t+1} = X_t + 1] + P[X_{t+1} = X_t - 1] = 1$$
- Example 1: Logistic function
 - $$P[X_{t+1} = X_t + 1] = \frac{1}{1 + \exp[-\alpha_1(L - X_t)]}$$

↑
level = long-term expectation
 - Mean reverting process
 - This is a Markov Process: Future independent of past given present
 - Markov Processes offer great computational benefit

• Modeling issue: State captures all relevant history

↳ Don't want too much information to float space

↳ State is sufficient history of future

• Example 2:

$$P[X_{t+1} = x_{t+1}] = \begin{cases} 0.5(1 - \alpha_2(x_t - x_{t-1})) & \text{if } t > 0 \\ 0.5 & \text{if } t = 0 \end{cases}$$

↳ To make this a Markov Process:

$$S_t = (x_t, x_t - x_{t-1})$$

• Example 3: Movement depends on entire history

$$P[X_{t+1} = x_{t+1}] = \begin{cases} \frac{1}{1 + (\frac{U_t + D_t}{D_t} - 1)^{\alpha_2}} & \text{if } t > 0 \\ 0.5 & \text{if } t = 0 \end{cases}$$

Unit Sigmoid Curve

↳ Movement bias to reverse entire history

$$S_t = (U_t, D_t)$$

↳ x_t is not in state, but can be derived from:

$$x_t = x_0 + U_t - D_t$$

- Make another example
- Generate Simulation Paths with different Strength parameters
- Run multiple simulation traces and plot distributions of outcomes at multiple points of time
 - ↳ Plot variances as a function of α
- Definition of Discrete Time, Countable States
 - Subset of state space denoted terminal space ends Markov process upon occurrence
 - Non-Stationary Transition probabilities have time-dependent transition probabilities.
 - In this course, if not specified, assume stationary Markov Process
 - ↳ Stationary Transition Probability Function
- $P(s, s') = P[s_{t+1} = s' | s_t = s]$ for $s \in N$
 $s' \in S$
- A lot of authors describe termination with absorbing state, but we do not.
 - ↳ Does not work well with MDP
 - ↳ Programmatically challenging
- In code we explicitly specify terminal states
- In design of software and model, consider:
 - ① Transition Probabilities, P
 - ② Probability distribution of start states, μ
- Episodic: All traces return to a terminating state

↳ Some processes do not terminate and are called Continuing.

- Code for MarkovProcess:

- Abstract methods must be defined in derived classes
- In code, specify types of each input and output

↳ Transition takes a state and returns a probability distribution over next states

→ Are programming concepts like generics, types, generators prerequisite for this course?

- isTerminal checks if transition distribution is None
- Generators create simulation trace in a lazy manner - only when somebody calls it.
 - ↳ Creates iterable step-by-step
 - ↳ Lets you deal with infinite states

- Instead of representing P as a matrix of transition probabilities, use *current* representation

Mapping $[s]$, Optional $[\text{FiniteDistribution}(s)]$
↑
State
Look at Distributions.py

- Look at printed values and graphical views of transition probabilities
- FiniteMarkovProcess is a concrete class of abstract class MarkovProcess
 - ↳ Makes concrete by implementing Transition method

- Inventory Example: Inventory Control

$$\begin{aligned} \alpha &:= \text{On-hand Inventory} \\ \beta &:= \text{On-order Inventory} \\ C &:= \text{Store Capacity} \end{aligned} \quad \left. \begin{array}{l} \\ \\ \end{array} \right\} \text{Constitutes State}$$

- Observe State $S_t = (\alpha, \beta)$ at 6 pm store close
- Order quantity := $\max(C - (\alpha + \beta), 0)$
- Order takes 36 hours to ship
- Demand is Poisson, i
- Inventory Sold = $\max(\alpha + \beta, i)$
- Why is it max and not min?

$$S := \{(x, y) : 0 \leq x + y \leq C\}$$

... See slide 17 and review textbook to frame Markov Process

- A Markov Chain w/o terminal state yields state frequency π

- For stationary process with finite states:

$\pi(s)$ is an eigenvector of P^T with eigenvalue of 1

- Markov Reward Process calculates expected accumulated reward, but does not optimize reward

$$P[(R_{t+1}, S_{t+1}) | S_t, \dots, S_0] = P[(R_{t+1}, S_{t+1}) | S_t]$$

→ MarkovRewardProcess in code is an abstract class

↳ transition-reward is new abstract method

↳ transition is implemented using transition-reward

* Course will focus on software design driven by math.

- Reward is often expressed as an expected reward (see slide 22)

↳ You can express reward function as a function of state alone

- You can extend inventory example with concept of reward

↳ There is a cost of "interest on inventory"

↳ There is also a cost associated with disappointed customers if inventory is not available ("soft money")

↳ Emotional factors for customers

$$P_a : N \times R \times S \rightarrow [0, 1]$$



$$N \rightarrow (S \times R \rightarrow [0, 1])$$

- Define Return G_t

↳ Enables convergence (enable mathematical tractability)

↳ Interest Rates, inflation, TVM

- If all segments terminate, set discount factor equal to 1

↳ E.g. $T=1$ for AI in car navigation

$$V(s) = \mathbb{E}[G_t | S_t = s]$$

↑
Value Function

- Bellman Equation is a recursive equation for $V(s)$

↳ See slide 27

↳ Requires matrix inversion, which can be intractable. Use DP instead!

- Intuition → Picture → Math → Programming